

```
In [58]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [59]: dataset=pd.read_csv('car data.csv')
```

```
In [60]: dataset.head(2)
```

```
Out[60]:
```

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	ritz	2014	5.59	27000	Petrol	Dealer	Manual	1
1	sx4	2013	9.54	43000	Diesel	Dealer	Manual	1

```
In [61]: dataset.isnull().sum()
```

```
Out[61]: Car_Name      0
Year      0
Present_Price  0
Kms_Driven  0
Fuel_Type    0
Seller_Type  0
Transmission  0
Owner        0
Selling_Price  0
dtype: int64
```

```
In [62]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Car_Name        301 non-null    object
1   Year            301 non-null    int64
2   Present_Price   301 non-null    float64
3   Kms_Driven      301 non-null    int64
4   Fuel_Type       301 non-null    object
5   Seller_Type     301 non-null    object
6   Transmission    301 non-null    object
7   Owner           301 non-null    int64
8   Selling_Price   301 non-null    float64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

## Car\_Name

```
In [63]: from sklearn.preprocessing import LabelEncoder
```

```
In [64]: LE= LabelEncoder()
```

In [65]: LE

Out[65]: LabelEncoder()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

In [66]: dataset['Car\_Name']=LE.fit\_transform(dataset['Car\_Name'])

In [67]: dataset.head()

Out[67]:

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	90	2014	5.59	27000	Petrol	Dealer	Manual	1
1	93	2013	9.54	43000	Diesel	Dealer	Manual	1
2	68	2017	9.85	6900	Petrol	Dealer	Manual	1
3	96	2011	4.15	5200	Petrol	Dealer	Manual	1
4	92	2014	6.87	42450	Diesel	Dealer	Manual	1

## Fuel\_Type

In [68]: dataset['Fuel\_Type'].unique()

Out[68]: array(['Petrol', 'Diesel', 'CNG'], dtype=object)

In [ ]:

In [69]: dataset['Fuel\_Type']=LE.fit\_transform(dataset['Fuel\_Type'])

In [70]: dataset.head(1)

Out[70]:

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	90	2014	5.59	27000	2	Dealer	Manual	1

## Seller\_Type

In [71]: dataset['Seller\_Type']=LE.fit\_transform(dataset['Seller\_Type'])

In [72]: dataset.head(1)

Out[72]:

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
0	90	2014	5.59	27000	2	0	Manual	1

## Transmission

In [73]: `dataset['Transmission']=LE.fit_transform(dataset['Transmission'])`

In [74]: `dataset.head(1)`

Out[74]:

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Own
0	90	2014	5.59	27000	2	0	1	

In [75]: `x=dataset.iloc[:, :-1]`

In [76]: `x`

Out[76]:

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Ow
0	90	2014	5.59	27000	2	0	1	
1	93	2013	9.54	43000	1	0	1	
2	68	2017	9.85	6900	2	0	1	
3	96	2011	4.15	5200	2	0	1	
4	92	2014	6.87	42450	1	0	1	
...	...	...	...	...	...	...	...	...
296	69	2016	11.60	33988	1	0	1	
297	66	2015	5.90	60000	2	0	1	
298	69	2009	11.00	87934	2	0	1	
299	69	2017	12.50	9000	1	0	1	
300	66	2016	5.90	5464	2	0	1	

301 rows × 8 columns

In [77]: `y=dataset['Selling_Price']`

In [78]: `y`

Out[78]:

0	3.35
1	4.75
2	7.25
3	2.85
4	4.60
...	
296	9.50
297	4.00
298	3.35
299	11.50
300	5.30

Name: Selling\_Price, Length: 301, dtype: float64

```
In [79]: from sklearn.preprocessing import StandardScaler
```

```
In [80]: SS=StandardScaler()
```

```
In [81]: SS
```

```
Out[81]: StandardScaler()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [82]: x=pd.DataFrame(SS.fit_transform(x), columns=x.columns)
```

```
In [83]: x.head()
```

```
Out[83]:
```

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	
0	1.074323	0.128897	-0.236215	-0.256224	0.500183	-0.737285	0.39148	-
1	1.191828	-0.217514	0.221505	0.155911	-1.852241	-0.737285	0.39148	-
2	0.212627	1.168129	0.257427	-0.773969	0.500183	-0.737285	0.39148	-
3	1.309332	-0.910335	-0.403079	-0.817758	0.500183	-0.737285	0.39148	-
4	1.152659	0.128897	-0.087890	0.141743	-1.852241	-0.737285	0.39148	-


```
In [84]: from sklearn.model_selection import train_test_split
```

```
In [85]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_sta
```

```
In [86]: from sklearn.linear_model import LinearRegression, Lasso, Ridge, ElasticNet
from sklearn.tree import DecisionTreeRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.svm import SVR
from sklearn.ensemble import RandomForestRegressor
```

```
In [87]: sns.heatmap(data=dataset.corr(),annot=True)
```

```
Out[87]: <Axes: >
```



## LinearRegression()

```
In [88]: LR=LinearRegression()
```

```
In [89]: LR
```

```
Out[89]: LinearRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [90]: LR.fit(x_train,y_train)
LR.score(x_train,y_train)*100 , LR.score(x_test,y_test)*100
```

```
Out[90]: (88.40630578239454, 84.65539666857805)
```

## Lasso

```
In [91]: LR1= Lasso(alpha=0.05)
LR1.fit(x_train,y_train)
LR1.score(x_train,y_train)*100 , LR1.score(x_test,y_test)*100
```

Out[91]: (88.35433202380113, 84.42023265451037)

## Ridge

```
In [92]: LR2= Ridge(alpha=10)
LR2.fit(x_train,y_train)
LR2.score(x_train,y_train)*100 , LR2.score(x_test,y_test)*100
```

Out[92]: (88.28628537091495, 84.16213595432282)

## ElasticNet

```
In [93]: LR3= ElasticNet(alpha=0.5)
LR3.fit(x_train,y_train)
LR3.score(x_train,y_train)*100 , LR3.score(x_test,y_test)*100
```

Out[93]: (84.00059239671332, 78.3177718663528)

## Decision Tree

```
In [94]: dt=DecisionTreeRegressor(max_depth=13)
dt.fit(x_train,y_train)
dt.score(x_train,y_train)*100 , dt.score(x_test,y_test)*100
```

Out[94]: (99.99983614119861, 95.47372436648206)

```
In [95]: from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```
In [96]: mean_squared_error(y_test,dt.predict(x_test)), mean_absolute_error(y_test,d
```

Out[96]: (1.042654049180328, 0.6585737704918032)

## Random Forest Regressor

```
In [97]: rf=RandomForestRegressor(n_estimators=100)
rf.fit(x_train,y_train)
rf.score(x_train,y_train)*100 , rf.score(x_test,y_test)*100
```

Out[97]: (98.30481566012408, 96.41847578267432)

```
In [98]: mean_squared_error(y_test,rf.predict(x_test)), mean_absolute_error(y_test,r
```

Out[98]: (0.8250250381967219, 0.5829622950819674)

## Support Vector regression

```
In [99]: SV=SVR()
SV.fit(x_train,y_train)
SV.score(x_train,y_train)*100 , SV.score(x_test,y_test)*100
```

Out[99]: (66.00840380338376, 78.48466914602926)

## KNeighborsRegressor

```
In [100]: Knn=KNeighborsRegressor()
Knn.fit(x_train,y_train)
Knn.score(x_train,y_train)*100 , Knn.score(x_test,y_test)*100
```

Out[100]: (91.06681012800678, 93.29996797075346)

In [ ]:

```
In [101]: rf.predict([[-1.275759,0.821718,-0.817924,-0.333500,0.500183,1.356327,-2.55

C:\Users\aditya\anaconda3\New folder\Lib\site-packages\sklearn\base.py:464:
UserWarning: X does not have valid feature names, but RandomForestRegressor was fitted with feature names
  warnings.warn(
```

Out[101]: array([0.4543])

```
In [102]: y_test
```

```
Out[102]: 177      0.35
289     10.11
228      4.95
198      0.15
60       6.95
...
234      5.50
296      9.50
281      2.10
285      7.40
182      0.30
Name: Selling_Price, Length: 61, dtype: float64
```

```
In [107]: new_data=pd.DataFrame([[ 'ritz',2014,5.59,27000,'Petrol','Dealer','Manual',0
new_data
```

```
Out[107]:
```

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner_Type
0	0	2014	5.59	27000	Petrol	Dealer	Manual	0

In [ ]:

```
In [109]: new_data['Fuel_Type']=LE.fit_transform(new_data['Fuel_Type'])
```


```
In [110]: new_data['Seller_Type']=LE.fit_transform(new_data['Seller_Type'])
```

```
In [111]: new_data['Transmission']=LE.fit_transform(new_data['Transmission'])
```

```
In [112]: new_data
```

```
Out[112]:
```

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner_Type
0	0	2014	5.59	27000	0	0	0	0




```
In [113]: new_data=pd.DataFrame(SS.transform(new_data), columns=new_data.columns)
```

```
In [114]: new_data
```

```
Out[114]:
```

	Car_Name	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner_Type
0	-2.4508	0.128897	-0.236215	-0.256224	-4.204665	-0.737285	-2.554408	-0.255408



```
In [115]: rf.predict(new_data)
```

```
Out[115]: array([3.4745])
```

```
In [ ]:
```

```
In [ ]:
```